

Synthetic Data Augmentation for Mispronunciation Detection in L2 English

Ella Hawkins

ella.hawkins@maine.edu

University of Southern Maine

Portland, Maine, USA

Abstract

Mispronunciation detection and diagnosis is an important part of computer-assisted language learning systems, where accurate feedback on pronunciation can help learners improve without the need for a human tutor. Most previous approaches to this problem rely on acoustic features extracted from speech audio, which can make them sensitive to noise and speaker variation. This paper explores a different approach: treating mispronunciation detection as a sequence labeling problem over International Phonetic Alphabet (IPA) transcriptions, rather than an acoustic classification task. Because annotated L2 transcription data is not readily available in large amounts, this paper also explores whether synthetically augmenting IPA sequences with errors can improve model performance. Three sequence models are trained and evaluated on various amounts of synthetically augmented data. Results show that a moderate amount of synthetic data (4,000 samples) improves phone- and word-level F1 for the BiLSTM and GRU, and that synthetic data at all evaluated scales significantly improves Transformer performance. Across all three models, synthetic augmentation reduces false positive error hallucinations, though it also tends to increase false negatives. These results suggest that IPA-based sequence modeling is a viable approach to mispronunciation detection, and that statistically-informed synthetic error generation can improve performance when real annotated data is scarce.

CCS Concepts

• Computing Methodologies → Natural Language Processing.

Keywords

Mispronunciation Detection, Data Augmentation, Phonetic Transcription

ACM Reference Format:

Ella Hawkins. 2018. Synthetic Data Augmentation for Mispronunciation Detection in L2 English. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 8 pages. <https://doi.org/XXXXXXX.XXXXXXX>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference acronym 'XX, Woodstock, NY

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-XXXX-X/2018/06

<https://doi.org/XXXXXXX.XXXXXXX>

1 Introduction

Mispronunciation detection and diagnosis (MDD) is an important component of computer-assisted language learning (CALL) systems, where it is used to provide learners with accurate feedback on their pronunciation. Most existing MDD systems approach this as an acoustic classification task, comparing a learner's speech against native pronunciation models using features extracted from the audio signal (e.g., phonetic embeddings [30] and phonological features [22]). While these approaches can be effective, they are sensitive to noise and speaker variation [7, 18], and they can struggle to distinguish sounds that are phonologically distinct but acoustically similar (i.e., they are different sounds linguistically, but they can sound similar) [22, 26].

An alternative approach is to treat mispronunciation detection as a problem in phonological sequences rather than acoustic signals. The International Phonetic Alphabet (IPA) represents speech sounds symbolically, with each symbol corresponding to one phone (sound). If a learner's speech can be transcribed into IPA, sequence models can operate directly on these representations without any access to the underlying audio. This allows phonological knowledge to be incorporated more directly, and avoids some of the issues that come with acoustic methods.

A key challenge with this approach is data scarcity, which has been identified in prior work on mispronunciation detection [3, 10, 28]. Annotated phoneme-level L2 transcriptions are difficult and expensive to obtain. To address this, we generate synthetic IPA transcriptions with errors injected based on statistics from a real Mandarin L1 learner dataset, and investigate whether training on this augmented data improves mispronunciation detection on real learner transcriptions.

To explore this approach, we train three sequence models (a BiLSTM, a GRU, and a Transformer) on IPA transcriptions from the speechocean762 dataset [32], with and without synthetic augmentation. Each model is evaluated at six training levels using phone-, word-, and sentence-level F1, and mean absolute error. By comparing models trained on real data only against models trained on both real and synthetic data at different scales, we can isolate the effect of synthetic data augmentation.

Our results show that training on a limited amount of synthetic data improves phone- and word-level F1 scores for the GRU and BiLSTM over their real-data-only baselines. Additionally, the transformer benefits from the synthetic data at all evaluated scales. Across all three models, the most consistent effect of synthetic data is fewer false positive error hallucinations; models trained only on real data frequently predict errors where none exist, while models trained with synthetic augmentation are more conservative. However, synthetic data models also tend to miss more errors.

These findings suggest that IPA-based sequence modeling is a viable approach to mispronunciation detection, and that synthetic error generation can be useful for addressing data scarcity.

The rest of the paper is organized as follows. Section 2 reviews prior work on mispronunciation detection, synthetic data augmentation, and the use of phonological knowledge in MDD systems. Section 3 describes our datasets, synthetic error generation procedure, models, evaluation metrics, and baselines. Section 4 presents experimental results and discussion. In section 5 we discuss limitations and future work, and in Section 6 we conclude the paper.

2 Related Work

Previous research on mispronunciation detection and diagnosis (MDD) has explored the following areas relevant to this work: automated pronunciation error detection, synthetic data augmentation for MDD and automatic speech recognition (ASR) models, and the use of phonological knowledge to model English L2 speech.

Automated pronunciation error detection. Approaches to automated MDD can be categorized by the type of input they rely on: acoustic alone, acoustic-phonetic, linguistic-based and hybrid [19]. Acoustic-only approaches are based solely on features extracted from the speech signal. Acoustic-phonetic methods augment these features with phoneme-level information, often using forced alignment to capture deviations from native pronunciations [12, 21]. Linguistic-based approaches incorporate knowledge from phonology, such as IPA transcriptions or segmental rules, either as input features or as training constraints [8, 16]. Finally, hybrid approaches combine acoustic, linguistic, and phonetic knowledge, often with neural networks or Transformer-based models, to directly model pronunciation errors and improve detection accuracy [5, 24, 27]. Most MDD systems rely heavily on acoustic representations, which makes them sensitive to noise and speaker variation, and limits their ability to distinguish phonologically distinct but acoustically similar sounds [3, 27].

Synthetic Data Augmentation. Data scarcity is a challenge in MDD, since collecting and annotating real learner speech is costly. Previous work addresses this by using synthetic or augmented data. It generally uses one of two approaches, the first of which is to modify speech audio directly. This approach generates synthetic speech using text-to-speech or accent conversion systems, often guided by phonological knowledge, and uses the resulting waveforms to augment ASR or MDD training data [1, 3, 4, 9]. These methods have shown that synthetically accented speech can improve performance on seen accents. The second approach is converting speech to phonemes, modifying the phonemes rather than the audio, then converting those to graphemes (written representations). This approach values phoneme-level modeling for understanding pronunciation variation but typically revisits acoustics later on, leaving open the question of whether symbolic phonological representations alone can support mispronunciation detection [16, 31].

Phonological Error Modeling. Similar work examines how phonological knowledge can be used to mimic L2 pronunciation patterns, and how this data can be used to guide models. Most high-performing acoustic models learn phonological patterns (e.g., substitutions, deletions) implicitly from training data [12, 27]. This approach is shown to be effective, however, it requires large amounts

of annotated data. Other approaches directly encode phonological insights and formalize those insights as rules to aid in data augmentation or to inform error analysis [16, 17, 19, 25, 29]. This approach is more data-efficient than the first approach, however, integration of the methods outlined by this approach into MDD systems remains largely unexplored.

To the best of our knowledge, no prior work has treated mispronunciation detection as a sequence labeling problem using IPA transcriptions, without access to acoustic data. Existing systems either operate only on audio, or use phonological representations as well as audio data. Similarly, prior work on synthetic data augmentation for MDD modifies audio, instead of speech transcriptions.

3 Methods

This section describes the datasets, error generation procedure, models, evaluation metrics, and baselines used in our experiments. We first describe the real learner data used for training and evaluation in Section 3.1, followed by the procedure for generating synthetic training data in Section 3.2. Section 3.3 describes the three sequence models used in our experiments, Section 3.4 presents the evaluation metrics, and Section 3.5 discusses the baseline systems.

3.1 speechocean762 Dataset

The speechocean762 dataset [32] is a collection of annotated L2 English transcriptions. It is designed for pronunciation assessment use. It contains 5,000 sentence transcriptions from 250 different speakers, Mandarin being the first language for all of them. Each transcription is annotated at the sentence, word, and phone level. Sentences are scored 0-10 for accuracy, fluency, completeness and prosody. Words are scored 0-10 for accuracy and stress of the word's pronunciation. At the phone level, each phone is scored 0-2 for pronunciation accuracy. For example, if a speaker pronounced the word "I" as "a", the word-level score would be 0. If they pronounced "thought" as "fought", the word-level score would be 7. Table 1 shows an annotated example sentence from the speechocean762 dataset. Each word is shown with its canonical (reference) IPA transcription and the learner's pronounced transcription. Words that differ from the canonical form are marked as errors and given mispronunciation metadata.

This word-level mispronunciation metadata, which is assigned to words that were pronounced with errors, records the canonical phone (the expected pronunciation), the pronounced phone (what the learner produced), and the phone's index in the word. A phone labeled with a pronounced value of "" indicates that the phone was deleted. Annotations are provided in the ARPAbet, which is a phonetic transcription system similar to the IPA; to represent phones, it uses sequences of English letters instead of IPA symbols. We converted all ARPAbet transcriptions to the IPA using the Python `phonocodes` library.

We split the speechocean762 dataset into training (80%), validation (10%), and test (10%) sets, stratified by sentence-level error count.

Figure 1 shows the score frequencies at both the word and sentence levels. The majority of words are scored as correct (score = 10), and most sentences are scored 8 or above, which shows that the majority of sentences and words were pronounced correctly.

Table 1: An example annotated sentence from the speechocean762 dataset. The “Pronounced” column shows the learner’s speech in IPA; errors are phones that differ from the reference transcription.

Word	Reference	Pronounced	Has Error
HE	h i	h i	No
JUST	d ʒ ʌ s t	d ʒ ʌ s t	No
TURNED	t ɛ ɹ n d	t ə ɹ z d	Yes (n → z)
AND	æ n d	æ n d	No
LEFT	l ɛ f t	l ɪ f t	Yes (ɛ → ɪ)




Figure 1: (a) Frequency of word-level pronunciation scores. (b) Frequency of sentence-level pronunciation scores.

Figure 2 shows the 10 most frequently mispronounced phones; it is worth noting that half of these, including the top 2, are vowels. Figure 3 shows the frequencies of the positions of errors in words, with the position rounded to the nearest tenth. Errors are most often near the end of the word, specifically, they occur on the last phone. A fair number of errors also occur near the middle of the word. These statistics are referenced in the probability distribution used to inject the synthetic errors.

3.2 Synthetic Data

Figure 4 summarizes our synthetic error injection pipeline. To generate synthetically augmented IPA transcriptions, we use the Jotschi/visual_genome-simple-en dataset on HuggingFace¹. This dataset contains simple English sentences produced by asking Mistral-7B-Instruct-v0.2-AWQ [6] to rephrase image captions from the Visual Genome dataset [11], which is a dataset of images from the Microsoft Common Objects in Context dataset [13], into a single sentence for a young child. Using LLM-generated sentences for children as a base is appropriate for this task because they are simple; the vocabulary and sentence structures are those commonly used




Figure 2: Top 10 frequently mispronounced phones.




Figure 3: Mispronunciation rate as a function of phone position within word.




Figure 4: Overview of synthetic error generation pipeline. Statistics are computed from the speechocean762 dataset and used to inject pronunciation errors into IPA transcriptions of simple English sentences.

in everyday spoken English, controlling the word- and structure-level complexities, which allows the focus to be on pronunciation patterns. For example, in Figure 4, we show the sentence “I was here”.

We convert these simple English sentences to the IPA as follows. Each sentence is first normalized to be fully capitalized, with

¹https://huggingface.co/datasets/Jotschi/visual_genome-simple-en

punctuation removed. Then, we convert each word to the ARPabet using the Carnegie Mellon Pronouncing Dictionary (CMUdict)². Words not found in the CMUdict are converted to the ARPabet by Python’s `g2p_en` library [20]. We do not use the G2P system for words already in the CMUdict to avoid inconsistencies. The ARPabet transcriptions are then converted to the IPA using the `phonocodes` Python library, the same conversion method used for the real learner data, ensuring consistency across both datasets. Figure 4 shows this conversion from the English sentence “I was here” to the IPA transcription: “ai wəz hiːr”.

To inject errors into these transcriptions, we first compute statistics from the train set of the `speechocean762` data. These statistics include the proportion of sentences with pronunciation errors, the average number of erroneous words in these sentences, phone error frequencies (see Figure 2), and the position of errors within words (see Figure 3). We then randomly select five sets of 2,000 sentences from the `visual_genome-simple-en` IPA transcriptions. We choose 2,000 sentences for the size of each synthetic set because this is 50% the size of the real (`speechocean762`) training data. Each sentence is probabilistically selected to contain errors based on the `speechocean762` statistics. For sentences selected, one or more words are randomly chosen, and one or more phones in each word are sampled to contain an error, again based on the `speechocean762` statistics. Figure 4 shows that two errors are injected into the transcription “ai wəz hiːr”, which makes it “ai wɛz hɛr” (this would be pronounced as “I wez her”).

3.3 Models

We use three sequence models (models that process ordered data): a gated recurrent unit (GRU), a bidirectional long short-term memory (BiLSTM), and a transformer. We use a shared vocabulary (of IPA tokens) across all models. We select these three architectures for the following reasons. Recurrent models such as the GRU and BiLSTM work well with sequential data with local context and have been widely applied to sequence labeling tasks in NLP [14, 15]. The GRU was chosen as a lighter alternative to the BiLSTM. It uses a single hidden state with reset and update gates, making it train efficiently while maintaining its performance.

The BiLSTM adds a cell state in addition to the hidden state, allowing it to model longer-range dependencies. We chose a BiLSTM rather than a unidirectional LSTM because phone-level error detection depends on context from both sides of a phone; whether a sound is an error can depend on what comes before and after it. We also exclude LSTM with attention as a standalone model because the Transformer, which is included, implements this idea more thoroughly by using multi-head self-attention over the entire sequence at once.

Together, the three architectures represent a progression from simple recurrent (GRU) to more complex recurrent (BiLSTM) to attention-based (Transformer), allowing us to compare how architecture affects performance under data scarcity.

For the BiLSTM and GRU, the two layers are stacked: the output of the first bidirectional layer (forward and backward hidden states concatenated, size $2 \times 256 = 512$) is fed as input to the second bidirectional layer. The final hidden state at each token position is

a 512-dimensional vector combining context from both directions. For the Transformer, the encoder consists of 4 stacked transformer encoder layers, 8 attention heads, and a feedforward dimension of 1,024.

All three models share the same classification heads: a phone head and a gap head. The phone head is a binary classifier applied at phone positions to predict whether each pronounced phone is correct or mispronounced. The gap head is a binary classifier, which, for each gap between phones, predicts whether a phone was deleted (i.e., if there should be a phone there, but the speaker did not produce it). Gap positions, which are slots between adjacent phones where a deletion may have occurred, are inserted explicitly into the input sequence during preprocessing, giving the model specific positions for deletion detection.

All three models use a learning rate of 3×10^{-4} , a dropout of 0.1, a batch size of 32, and train for 20 epochs, selecting the best checkpoint based on phone-level F1 on the validation set.

3.4 Evaluation Metrics

We evaluate each model at the phone, word, and sentence levels. The phone-level F1 is computed from binary error classification of each phone (correct = 0, error = 1). The word- and sentence-level F1 scores are computed similarly, treating any word with at least one phone error as an erroneous word, and each sentence with at least one phone error as an erroneous sentence. Additionally, we report the mean absolute error (MAE) between the predicted and ground-truth count of errors in each sentence.

Statistical significance is assessed for each training level against the real-only baseline using approximate randomization ($n = 10,000$). Training levels that performed significantly better (with $p < 0.05$) than the baseline on a specific metric are indicated by an asterisk.

3.5 Baseline Systems

To study the effect of synthetically augmented data on three different model types, we create a baseline for each model type. The baselines use the same architecture as the models trained with synthetic data, but they are trained only on the real learner data (`speechocean762`). Using the same architecture and evaluation metrics ensures a fair comparison, as all other experimental conditions remain identical, allowing us to isolate the effect of synthetic data augmentation.

The code for this project is publicly available.³

4 Experimental Results

We evaluate each model at six training levels: the real-only baseline (4,000 samples from the `speechocean762` dataset [32]), and five conditions that use 2,000, 4,000, 6,000, 8,000, and 10,000 samples of synthetic data, in addition to the real data.

GRU Results. As shown in Table 2, the GRU trained on 4,000 synthetic samples performed better than the baseline GRU in all categories besides sentence-level F1. This is because the baseline tended to predict errors in sentences that actually had errors, but it often predicted them in the wrong spot. This caused the sentence-level correct prediction count to be higher than the phone- and

² <https://github.com/cmuspinx/cmu dict>

³ <https://github.com/ella84211/synthetic-mispronunciation-data>

Table 2: Evaluation results for GRU. All models trained on 4,000 real samples, in addition to the indicated amount of synthetic data. * indicates performance significantly better than the baseline.

Synthetic Samples	Phone F1	Word F1	Sentence F1	MAE
0 (baseline: 4,000 real)	0.3917	0.3827	0.5185	0.5920
2,000 synthetic	0.4122	0.3136	0.4516	0.5580
4,000 synthetic	0.4601*	0.3988	0.5029	0.5380
6,000 synthetic	0.2072	0.2222	0.3583	1.7920
8,000 synthetic	0.2079	0.2188	0.3594	1.8660
10,000 synthetic	0.2095	0.2247	0.3683	1.7700

Table 3: Evaluation results for BiLSTM. All models trained on 4,000 real samples, in addition to the indicated amount of synthetic data. * indicates performance significantly better than the baseline.

Synthetic Samples	Phone F1	Word F1	Sentence F1	MAE
0 (baseline: 4,000 real)	0.2806	0.2945	0.4541	1.0640
2,000 synthetic	0.2833	0.2864	0.4241	1.0720
4,000 synthetic	0.4323*	0.3940*	0.4744	0.5640*
6,000 synthetic	0.1997	0.2158	0.3514	1.7080
8,000 synthetic	0.1665	0.1837	0.3345	2.2960
10,000 synthetic	0.1512	0.1625	0.3316	2.7700

word-level correct prediction counts. For instance, in the sentence, “he was face to face with the girl”, errors occurred in the first word (“he” was pronounced as “she”) and the last word (“girl” was pronounced as “curl”). The baseline predicted only one error in the sentence, in the sixth word. This means that it correctly predicted the sentence as erroneous, but did not classify the phones and words correctly. Therefore, the sentence-level F1 score was higher than the phone- and word-level F1 scores.

The most frequently occurring pattern was that the GRU trained only on real data usually over-predicts errors, compared to the GRUs trained on synthetic data. For example, in the sentence “how long was he going to be here”, which was pronounced “HOE long was he going to be HER”, the GRU trained on 4,000 samples of synthetic data correctly predicted one error in the first word and two errors in the second, however, the GRU trained on real data predicted three extra errors between two extra words, neither of which actually had errors. This same pattern is observed in the sentence “there are going to be more”, where “there” was pronounced like “thur”. While the GRU model trained on the synthetic data correctly predicts the two errors in the sentence, both of which appear in the first word, the GRU model trained only on the real data predicts three errors in the first and third words, and one error in the second word.

The GRU’s tendency to over-predict errors when trained only on real data likely reflects the limited error diversity of the training data. With only 4,000 real samples, the model does not see enough examples of correctly-pronounced sentences to reliably distinguish them from erroneous ones. The addition of synthetic data, the majority of which represents correctly-pronounced phones, expands the model’s exposure to the “correct” class and influences its decision boundary.

BiLSTM Results. Table 3 shows that training the BiLSTM on 4,000 samples of synthetic data significantly improves the MAE and the phone- and word- level F1 scores, compared to the real-only baseline. Training on 2,000 samples of synthetic data appears to have little effect on the results, and performance decreases when training on 6,000 synthetic samples or more.

Many of the instances where this BiLSTM performed better than the baseline occurred when the synthetic data BiLSTM correctly predicted a small amount of errors, but the baseline hallucinated a large number of errors. This pattern is also observed with the GRU models. For example, in the sentence “these were all little girls”, the word “these” was pronounced as “there”. The synthetic data BiLSTM correctly predicted this error. However, the BiLSTM trained only on real data predicted that there were two errors in the first word, and one error in each of the next two words, hallucinating a total of three errors. Another instance of this pattern is in the sentence “how long was he going to be here”; the synthetic data BiLSTM correctly predicted one error in the first word and two errors in the second. However, the baseline predicted two errors in the first word, one error in the second, two errors in the fifth word, and two errors in the last word.

On the other hand, the BiLSTM trained on 4,000 synthetic samples tends to under-predict errors. For example, in the sentence, Linda got the sunglasses”, the word “got” was pronounced as “goat”. The synthetic data BiLSTM did not predict any errors, however, the baseline correctly predicted this error. A similar pattern appeared in the sentence, “there it was right next to him”. “There” was pronounced similar to “cha”, however, the synthetic data BiLSTM predicted the sentence as correct. The real data BiLSTM correctly predicted these three errors.

While the BiLSTM trained on 4,000 synthetic samples performs better than the real-only baseline, the BiLSTMs trained on 6,000, 8,000, and 10,000 synthetic samples tend to over-predict errors. For instance, in the sentence, “she wants to ask you herself”, which did not have any errors, the BiLSTMs trained on more than 4,000 synthetic samples each predicted 3-5 errors, whereas the one trained on 4,000 synthetic samples correctly predicted none.

The BiLSTM’s behavior at 4,000 synthetic samples, reduced false positives but increased false negatives, suggests that the additional data helped the model learn when not to predict an error, but the synthetic error distribution did not fully capture the diversity of real learner errors, leaving the model less able to predict rare patterns. The degradation at 6,000+ synthetic samples suggests that beyond a certain point, the synthetic data begins to dilute the real learner signal rather than strengthen it, making the model predict patterns that are more common in the synthetic distribution.

Transformer Results. The baseline transformer, trained on only real data, performs very poorly. It predicts errors in every sample, sometimes predicting whole words and sentences as erroneous. This is incorrect, as only about 20% of the sentences contain errors, of the 500 sentences in the test set. Additionally, of the sentences that contain errors, it is rare that full words and sentences are erroneous.

Table 4 shows that training the transformer on synthetic data at all evaluated scales significantly improves the performance compared to the baseline. Within the observed range (up to 10,000 samples), as the amount of synthetic data increases, the frequency

Table 4: Evaluation results for transformer. All models trained on 4,000 real samples, in addition to the indicated amount of synthetic data. * indicates performance significantly better than the baseline.

Synthetic Samples	Phone F1	Word F1	Sentence F1	MAE
0 (baseline: 4,000 real)	0.0695	0.1192	0.3221	4.7140
2,000 synthetic	0.1909*	0.2016*	0.3834*	1.7980*
4,000 synthetic	0.2298*	0.2976*	0.4015*	0.9680*
6,000 synthetic	0.1093*	0.1402*	0.3221	3.5520*
8,000 synthetic	0.2804*	0.3435*	0.4635*	0.7760*
10,000 synthetic	0.3048*	0.3437*	0.4434*	0.7840*

of hallucinated errors decreases. For example, the sentence “I will put pressure on him and make him miss”, was pronounced completely correctly. While the transformer trained on 10,000 synthetic samples correctly predicted that the sentence had no errors, the baseline predicted errors in every word besides the first, for a total of twenty-one errors, when none were pronounced.

Another example of this over-prediction is in the sentence, “it will have a body again”, where “body” was pronounced as “buddy”. The transformer trained on 10,000 synthetic samples correctly predicted that the word had an error, but it predicted that this error was at index 3, instead of index 1. However, this model performed better than the baseline, which predicted an error in the third word, and two errors in the fifth word. This shows that even when the baseline does not classify the majority (or all) of a sentence as erroneous, in most cases, it still predicts too many errors, whereas the model trained on synthetic data rarely has this issue.

The baseline transformer’s severe over-prediction is consistent with well-known behavior of transformers trained on small datasets: without sufficient data, the self-attention mechanism cannot learn meaningful patterns, and the model defaults to predicting the majority class [2, 23]. Since the training loss penalizes missed errors (which are more rare, but still important), the model learns to predict errors everywhere to minimize this penalty. The transformer’s consistent improvement with more synthetic data reflects the architecture’s general data hunger; unlike recurrent models, which build in biases about sequential order, the transformer must learn positional and sequential structure entirely from data.

Global Analysis. Two patterns were observed across all three models. First, training on the synthetic data helps all models prevent error hallucination; the models trained on both synthetic and real data consistently hallucinated fewer errors than the models trained on only real data. For the BiLSTM and the GRU, this is likely because the synthetic data allowed them to see more instances of errors in the synthetic data than what appeared in the real data. For the transformer, the synthetic data greatly improved the performance. As described above, there were multiple instances where the baseline transformer predicted every phone in a word as an error, for multiple words in a sentence, when there were no errors in the sentence at all. The transformer trained on synthetic data was able to avoid this problem. This is likely because of the transformer architecture; transformers often perform better when trained on larger amounts of data.

Second, while the models trained on synthetic data seem to hallucinate less, they tend to become more conservative when predicting

Table 5: Example model predictions for one word. 1 = error. B = baseline (real data only) and T = best synthetic condition (4,000 samples for GRU and BiLSTM; 10,000 samples for the transformer).

Phones	k	l	au	d
Ground Truth	0	0	0	0
GRU B	1	1	0	0
GRU T	0	0	0	0
BiLSTM B	0	1	0	1
BiLSTM T	0	0	0	0
Transformer B	1	1	1	1
Transformer T	0	0	0	0

Table 6: Sentence-level false positive (FP) and false negative (FN) error counts by model. B = baseline (real data only) and T = best synthetic condition (4,000 samples for GRU and BiLSTM; 10,000 samples for the transformer).

Model	FP	FN	Total Errors
GRU B	190	12	202
GRU T	35	52	87
BiLSTM B	90	12	102
BiLSTM T	68	45	113
Transformer B	404	0	404
Transformer T	69	49	118

errors. These models frequently under-predict errors, often failing to identify any errors in a sentence that contains several. A possible cause for this is that the synthetic data reduces the models’ abilities to detect rare error patterns. The synthetic data BiLSTMs and GRUs often under-predict so much that they incorrectly classify whole words, predicting them as correct, even when every phone in the word is pronounced with an error. This pattern continues with the transformers, with the baseline transformer predicting too many errors, often at incorrect positions, and many of the synthetic data transformers not predicting any errors. A possible cause for this is that the synthetic data reduces the models’ abilities to detect rare error patterns, as the errors in the synthetic data are injected based on probability. This means that the majority of the errors in the synthetic data are more frequent errors; therefore, the model may not see a large range of uncommon patterns.

Table 5 shows example predictions for each model type (the real data baseline and the best-performing model trained on synthetic data) on the word “cloud” (IPA: k l au d), which was pronounced correctly (with no errors). “0” indicates that the model predicted the phone to be correct, and “1” indicates that the model predicted the phone to be an error. The transformer baseline does the worst; it predicts every phone to be an error. The GRU and BiLSTM baselines also hallucinate errors that do not exist. However, all three of the models trained on synthetic data correctly predict that the word was pronounced correctly.

Error Analysis. Errors in the models’ predictions are classified as either false positive or false negative. A false positive error means

that the model predicted an error where there were none, and a false negative error means that the model did not predict an error where one existed. For each model type, we consider the baseline (trained only on real data) and the best-performing model of that type (trained on synthetic data).

As shown in Table 6, for all models, with the exception of the GRU trained on 4,000 samples of synthetic data, the majority of the models' prediction errors were false positives. The baseline transformer model predicts the most errors, of the models evaluated. Its prediction was incorrect for 404 sentences. All of these predictions were false positives. This shows that the baseline transformer almost always predicts that errors exist, even when there are none.

Similarly, the BiLSTM and GRU baselines also show a tendency towards false positives, but they predict errors less aggressively. They each predicted a number of sentences to have no errors, unlike the transformer, which predicted that most sentences had errors. The predictions from both the BiLSTM and the GRU had a fair number of both false positive and false negative errors, but they both still had far more false positives than false negatives.

Training with synthetic data greatly changes this error distribution. For the BiLSTM trained on 4,000 samples of synthetic data, the false positive errors were reduced from 90 (from the BiLSTM trained only on real data) to 68. However, the false negative errors increased from 12 to 45. The GRU follows a similar pattern; the false positives were reduced from 190 to 35, but the false negatives were increased from 12 to 52.

5 Discussion

The high number of false positives from the baseline models' predictions indicates that training on limited real data makes it difficult for models to detect when a sentence is pronounced correctly. This suggests that additional data is necessary for the models to better learn pronunciation error patterns. While synthetic data does improve this issue, the increase in false negatives across all three models trained on synthetic data suggests that the synthetic data may not realistically represent more uncommon error patterns. Overall, these results suggest that synthetic augmentation does improve mispronunciation detection performance, though the effect varies by model and scale.

One limitation of our work is that the synthetic data does not represent more uncommon errors or error context. The errors are injected into the synthetic data based on the frequencies that they appear in the speechocean762 data. There are certain errors (such as removing an "s") that occur more often before or after a specific phone. For example, the "s" is deleted more often after consonants than vowels. In this study, many models are not learning the context because it is not in the synthetic data; including the context could help them have more accurate predictions. Future work could include oversampling more uncommon errors and having a more complex system that considers context, or developing a rule-based approach using linguistic knowledge for synthetic error injection.

Additionally, the synthetic data does not imitate error clusters in words and sentences. Therefore, the models are learning that errors are somewhat evenly distributed, whereas in real learner speech, errors are likely to be clustered, with many errors appearing close to each other in both words and sentences. For example,

in the word "strengths" (IPA: st.ɪŋkθs), which has eight phones, learners may mispronounce the whole consonant cluster at the end. However, our current synthetic error injection approach treats phones independently, making such clusters unlikely. Future work could include studying the attributes of these clusters and developing a system to make the appearance of the synthetic errors more realistic.

Another limitation is that we did not perform hyperparameter ablations on the model architectures. For example, varying the model dimension (e.g., 64, 128, or 256 for the transformer) or the number of layers (for all models) would likely affect performance. It is possible that smaller architectures would be less prone to overfitting on the real-only training data and would therefore require less synthetic augmentation to achieve good performance. Future work could analyze this.

6 Conclusion

This paper explored mispronunciation detection as a sequence labeling problem over IPA transcriptions, using synthetic data augmentation to address the scarcity of annotated L2 learner data. We trained a GRU, BiLSTM, and Transformer across six training conditions varying the amount of synthetic data. For the GRU and BiLSTM, 4,000 synthetic samples produced the strongest results, with statistically significant improvements in phone- and word-level F1. The Transformer improved at all evaluated scales, generally benefiting from more data. Across all models, synthetic augmentation consistently reduced false positive hallucinations, though at the cost of more false negatives — likely because the synthetic data overrepresents common error patterns and underrepresents rare ones. Future work should focus on improving synthetic error realism by incorporating error context and clustering, and on analyzing model performance with different hyperparameters.

References

- [1] Martijn Bartelds, Nay San, Bradley McDonnell, Dan Jurafsky, and Martijn Wiering. 2023. Making More of Little Data: Improving Low-Resource Automatic Speech Recognition Using Data Augmentation. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 715–729. doi:10.18653/v1/2023.acl-long.42
- [2] Maria Cristina De Cola, Fabiana Pia Vitiello, Angelo Quartarone, and Maria Lui. 2026. Can Machine Learning be Effective with Small Datasets? Insights from Modeling in Neuroscience. In *Health Information Science*, Emil R. Kaburuan and Sam Goundar (Eds.). Springer Nature Singapore, Singapore, 214–225.
- [3] Yassine EL Kheir, Shammur Chowdhury, Ahmed Ali, Hamdy Mubarak, and Shazia Afzal. 2023. SpeechBlender: Speech Augmentation Framework for Mispronunciation Data Generation. In *9th Workshop on Speech and Language Technology in Education (SLaTE)*. 26–30. doi:10.21437/SLaTE.2023-6
- [4] Lev Finkelstein, Heiga Zen, Norman Casagrande, Chun an Chan, Ye Jia, Tom Kenter, Alex Petelin, Jonathan Shen, Vincent Wan, Yu Zhang, Yonghui Wu, and Robert Clark. 2022. Training Text-To-Speech Systems From Synthetic Data: A Practical Approach For Accent Transfer Tasks. In *Interspeech 2022*. 4571–4575. doi:10.21437/Interspeech.2022-10115
- [5] Kaiqi Fu, Jones Lin, Dengfeng Ke, Yanlu Xie, Jinsong Zhang, and Binghui Lin. 2021. A Full Text-Dependent End to End Mispronunciation Detection and Diagnosis with Easy Data Augmentation Techniques. *arXiv e-prints*, Article arXiv:2104.08428 (April 2021), arXiv:2104.08428 pages. arXiv:2104.08428 [cs.CL] doi:10.48550/arXiv.2104.08428
- [6] Albert Q. Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, L  lio Renard Lavaud, Marie-Anne Lachaux, Pierre Stock, Teven Le Scao, Thibaut Lavril, Thomas Wang, Timoth  e Lacroix, and William El Sayed. 2023. Mistral 7B. arXiv:2310.06825 [cs.CL] <https://arxiv.org/abs/2310.06825>

- [7] T. Tuithung Kapang Legoh, Utpal Bhattacharjee. 2015. Features and Model Adaptation Techniques for Robust Speech Recognition: A Review. *Communications on Applied Electronics* 1, 2 (January 2015), 18–31. doi:10.5120/cae-1507
- [8] Fedor Kitashov, Elizaveta Svitanko, and Debojyoti Dutta. 2018. Foreign English Accent Adjustment by Learning Phonetic Patterns. arXiv:1807.03625 [cs.SD] <https://arxiv.org/abs/1807.03625>
- [9] Philipp Klumpp, Pooja Chitkara, Leda Sari, Prashant Serai, Jilong Wu, Irina-Elena Veliche, Rongqing Huang, and Qing He. 2023. Synthetic Cross-accent Data Augmentation for Automatic Speech Recognition. *arXiv e-prints*, Article arXiv:2303.00802 (March 2023), arXiv:2303.00802 pages. arXiv:2303.00802 [cs.CL] doi:10.48550/arXiv.2303.00802
- [10] Daniel Korzekwa, Jaime Lorenzo-Trueba, Thomas Drugman, Shira Calamaro, and Bozena Kostek. 2021. Weakly-Supervised Word-Level Pronunciation Error Detection in Non-Native English Speech. In *Interspeech 2021*. 4408–4412. doi:10.21437/Interspeech.2021-38
- [11] Ranjay Krishna, Yuke Zhu, Oliver Groth, Justin Johnson, Kenji Hata, Joshua Kravitz, Stephanie Chen, Yannis Kalantidis, Li-Jia Li, David A. Shamma, Michael S. Bernstein, and Li Fei-Fei. 2017. Visual Genome: Connecting Language and Vision Using Crowdsourced Dense Image Annotations. *Int. J. Comput. Vision* 123, 1 (May 2017), 32–73. doi:10.1007/s11263-016-0981-7
- [12] Kun Li, Xiaojun Qian, Helen Meng, Kun Li, Xiaojun Qian, and Helen Meng. 2017. Mispronunciation Detection and Diagnosis in L2 English Speech Using Multidistribution Deep Neural Networks. *IEEE/ACM Trans. Audio, Speech and Lang. Proc.* 25, 1 (Jan. 2017), 193–207. doi:10.1109/TASLP.2016.2621675
- [13] Tsung-Yi Lin, Michael Maire, Serge Belongie, Lubomir Bourdev, Ross Girshick, James Hays, Pietro Perona, Deva Ramanan, C. Lawrence Zitnick, and Piotr Dollár. 2015. Microsoft COCO: Common Objects in Context. arXiv:1405.0312 [cs.CV] <https://arxiv.org/abs/1405.0312>
- [14] Di Liu and Xinyi Zou. 2018. Sequence Labeling of Chinese Text Based on Bidirectional Gru-Cnn-Crf Model. In *2018 15th International Computer Conference on Wavelet Active Media Technology and Information Processing (ICCWAMTIP)*. 31–34. doi:10.1109/ICCWAMTIP.2018.8632570
- [15] Xuezhe Ma and Eduard Hovy. 2016. End-to-end Sequence Labeling via Bi-directional LSTM-CNNs-CRF. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Katrin Erk and Noah A. Smith (Eds.). Association for Computational Linguistics, Berlin, Germany, 1064–1074. doi:10.18653/v1/P16-1101
- [16] Margot Masson and Julie Carson-berndsen. 2023. Investigating Phoneme Similarity with Artificially Accented Speech. In *Proceedings of the 20th SIGMORPHON workshop on Computational Research in Phonetics, Phonology, and Morphology*, Garrett Nicolai, Eleanor Chodroff, Frederic Mailhot, and Çağrı Çöltekin (Eds.). Association for Computational Linguistics, Toronto, Canada, 49–57. doi:10.18653/v1/2023.sigmorphon-1.6
- [17] Margot Masson and Julie Carson-Berndsen. 2024. Investigating the Use of Synthetic Speech Data for the Analysis of Spanish-Accented English Pronunciation Patterns in ASR. In *Synthetic Data's Transformative Role in Foundational Speech Models*. 81–85. doi:10.21437/SynData4GenAI.2024-17
- [18] Ji Ming, Timothy J. Hazen, James R. Glass, and Douglas A. Reynolds. 2007. Robust Speaker Recognition in Noisy Conditions. *IEEE Transactions on Audio, Speech, and Language Processing* 15, 5 (2007), 1711–1723. doi:10.1109/TASL.2007.899278
- [19] Lakshani Nissanka, Banuka Athuraliya, and Sahana Priyanayana. 2025. Systematic review of feature-based approaches to mispronunciation detection. *The JALT CALL Journal* 21, 3 (Dec. 2025), 102732. doi:10.29140/jaltcall.v21n3.102732
- [20] Jongseok Park, Kyubyong Kim. 2019. g2pE. <https://github.com/Kyubyong/g2pE>
- [21] Linkai Peng, Yingming Gao, Binghui Lin, Dengfeng Ke, Yanlu Xie, and Jinsong Zhang. 2022. Text-Aware End-to-end Mispronunciation Detection and Diagnosis. arXiv:2206.07289 [cs.SD] <https://arxiv.org/abs/2206.07289>
- [22] Mostafa Shahin, Julien Epps, and Beena Ahmed. 2025. Phonological level wav2vec2-based Mispronunciation Detection and Diagnosis method. *Speech Communication* 173 (2025), 103249. doi:10.1016/j.specom.2025.103249
- [23] Ran Shao and Xiao-Jun Bi. 2022. Transformers Meet Small Datasets. *IEEE Access* 10 (2022), 118454–118464. doi:10.1109/ACCESS.2022.3221138
- [24] Qin Shi, Kun Li, ShiLei Zhang, Stephen M. Chu, Ji Xiao, and Zhijian Ou. 2010. Spoken English assessment system for non-native speakers using acoustic and prosodic features. In *Interspeech 2010*. 1874–1877. doi:10.21437/Interspeech.2010-541
- [25] Yui Suzukida and Kazuya Saito. 2022. What is second language pronunciation proficiency? An empirical study. *System* 106 (02 2022), 102754. doi:10.1016/j.system.2022.102754
- [26] Si Wei, Guoping Hu, Yu Hu, and Ren-Hua Wang. 2009. A new method for mispronunciation detection using Support Vector Machine based on Pronunciation Space Models. *Speech Communication* 51, 10 (2009), 896–905. doi:10.1016/j.specom.2009.03.004 Spoken Language Technology for Education.
- [27] Minglin Wu, Kun Li, Wai-Kim Leung, and Helen Meng. 2021. Transformer Based End-to-End Mispronunciation Detection and Diagnosis. In *Interspeech 2021*. 3954–3958. doi:10.21437/Interspeech.2021-1467
- [28] Mu Yang, Kevin Hirschi, Stephen Daniel Looney, Okim Kang, and John H.L. Hansen. 2022. Improving Mispronunciation Detection with Wav2vec2-based Momentum Pseudo-Labeling for Accentedness and Intelligibility Assessment. In *Interspeech 2022*. 4481–4485. doi:10.21437/Interspeech.2022-11039
- [29] Shanting Yang, Manting Wu, Piaoyan Jiang, Xin Zhou, and Hongyan Wang. 2023. The Mutual Intelligibility of English Front Vowels by Cantonese, Mandarin and English Native Speakers. In *Proceedings of the 2023 8th International Conference on Big Data and Computing (Shenzhen, China) (ICBDC '23)*. Association for Computing Machinery, New York, NY, USA, 55–60. doi:10.1145/3624288.3624296
- [30] Wenxuan Ye, Shaoguang Mao, Frank Soong, Wenshan Wu, Yan Xia, Jonathan Tien, and Zhiyong Wu. 2021. An Approach to Mispronunciation Detection and Diagnosis with Acoustic, Phonetic and Linguistic (APL) Embeddings. In *ICASSP 2022*.
- [31] Saierdaer Yusuyin, Te Ma, Hao Huang, Wenbo Zhao, and Zhijian Ou. 2025. Whistle: Data-Efficient Multilingual and Crosslingual Speech Recognition via Weakly Phonetic Supervision. *IEEE Transactions on Audio, Speech and Language Processing* 33 (2025), 1440–1453. doi:10.1109/TASLP.2025.3550683
- [32] Junbo Zhang, Zhiwen Zhang, Yongqing Wang, Zhiyong Yan, Qiong Song, Yukai Huang, Ke Li, Daniel Povey, and Yujun Wang. 2021. speechocan762: An Open-Source Non-Native English Speech Corpus for Pronunciation Assessment. In *Interspeech 2021*. 3710–3714. doi:10.21437/Interspeech.2021-1259

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009